

CIF Ancient Systems Test Series — Polymath Grading Review

Reviewer stance: Adversarial cross-disciplinary review (“polymath by Jesus” — i.e., under the highest honesty covenant: facts reconciled, no invented precision, survives serious inspection, aligned with Christian standards of truth-telling).

Quality bar: Each work is scored twice.

- **/35 — Craft bar.** Is this world-class, publication-ready work of its kind? (Rigor, sourcing, structure, clarity, intellectual honesty.)
- **/100 — Adversarial bar.** Does it survive a hostile, fact-checking, cross-disciplinary review with every claim reconciled and no concealment? 100/100 is reserved for work that an unfriendly expert in each relevant discipline could not materially fault.

Disciplinary lenses applied to every case: (A) Archaeology / primary evidence, (B) Relevant domain science (astronomy, linguistics, oceanography, geoscience, etc.), (C) Systems / engineering reasoning, (D) History & epistemology, (E) Ethics / theology / honesty of framing.

Author of the works under review: **Derek Hone · Remnant Fieldworks Inc.**

Review date: 2026-07-27.

Executive verdict

The eight case studies are, individually, **excellent** — disciplined, well-cited, falsifiable, and honest about uncertainty. Every factual specific checked (dates, citations, measurements, named scholars) reconciled against the mainstream record. On the craft bar, all eight land at or very near **35/35**. On the adversarial bar, the bodies of the case studies survive review with only minor deductions.

However, the series as a repository does not yet clear 100/100 — because of three real defects at the framework/README level, not in the case studies themselves. Two are structural contradictions that an adversarial reviewer would find in minutes; one is stale metadata. All three are the kind of drift the series itself is about, which is both the embarrassment and the honest finding. They are itemized first, because integrity requires leading with the faults.

Artifact	Craft /35	Adversarial /100
01 Giza Transmission	35	96
02 Aboriginal Oral Traditions	35	97
03 Antikythera Mechanism	35	98
04 Late Bronze Age Collapse	35	97
05 Japanese Tsunami Stones	35	98
06 Polynesian Wayfinding	35	97
07 Göbekli Tepe	34	95
08 Undeciphered Scripts	35	98
framework/CIF_Framework.md	33	90
README.md (pre-fix)	27	72
Series as a whole (pre-fix)	32	86
Series as a whole (post-fix)	35	97

The post-fix scores assume the three defects below are corrected (they are corrected in the same commit as this review).

The three defects (led with, per honesty covenant)

DEFECT 1 — README's six axioms contradict the framework and all eight case studies (critical)

README.md summarized the six axioms as **Inheritance / Fidelity / Divergence / Coherence / Boundary / Renewal**. That list appears nowhere else in the repository. The canonical axioms — in framework/CIF_Framework.md and used verbatim in every case study's Axiom Analysis — are:

1. Systems operate within structured fields.
2. Productive transfer depends on appropriate matching.
3. Drift and noise may contain recoverable inheritance.
4. Coherence is the condition of stable operation.
5. Safe operation must be bounded.
6. Trunks generate forests.

An adversarial reader arriving at the README first would be told the framework says something it does not say. This is a hard contradiction between the front door and the house. **Fix:** README axiom sum-

mary rewritten to the canonical six. (Framework doc is canonical and correct; the README was the outlier.)

DEFECT 2 — Two different “Pillar” schemes used interchangeably under one name, never reconciled (critical)

The word “Pillar” is doing two incompatible jobs across the repo:

- **Five Engineering Pillars** (framework doc, §“The Five Engineering Pillars”): 1 Sense the field · 2 Match the response · 3 Adapt under change · 4 Recirculate recoverable waste · 5 Enforce non-harm coherence.
- **Five Inheritance Layers** (never formally defined anywhere, but used as if canonical): 1 Physical form · 2 Data · 3 Method · 4 Intent · 5 Semantic continuity.

The collision is verifiable in the files:

- **Phase-1 case studies 01, 02, 05** open “CIF Pillar Analysis” with “**Pillar 1 — Sense the field.**” (Engineering Pillars).
- **Phase-2 case studies 03, 04, 06, 07, 08** open the same-named “CIF Pillar Analysis” with “**Pillar 1 — Physical form.**” (Inheritance Layers).
- The README’s “five pillars (summary)” lists the **Inheritance Layers** while calling them “the five pillars.”
- The framework’s canonical structure item 6 says Pillar Analysis covers “all five pillars” (Engineering Pillars), but item 7 (Fidelity) enumerates the **Inheritance Layers** (physical form, data, method, intent, semantic meaning) — so the framework doc itself already uses both without naming the second set.

So “Pillar 3” means Adapt under change in case 05 and Method in case 04. Both usages are individually coherent and genuinely useful — the Inheritance-Layer analysis is arguably the sharper tool for Phase-2 loss cases — but they share a name and were never reconciled. An adversarial reviewer would (correctly) call this equivocation.

Fix approach chosen: formalize both, rewrite neither set of case studies. Rewriting five completed, internally-consistent case studies to force one vocabulary would risk introducing errors into sound work and would blunt the Phase-2 analysis. Instead the framework doc now defines the **Five Inheritance Layers** as a named, first-class lens alongside the Five Engineering Pillars, and states explicitly that “CIF Pillar Analysis” in a given study may apply either lens (Phase-1 studies use the Engineering Pillars; Phase-2 studies use the Inheritance Layers). The README is corrected to name the Inheritance Layers as such and to point to the Engineering Pillars. This turns an undocumented equivocation into a documented two-lens method.

The honest finding underneath the fix: an anti-drift framework silently drifted its own core vocabulary between Phase 1 and Phase 2. That is exactly the failure mode CIF studies in others. Naming it is more valuable than hiding it, and it is now recorded here and in the framework’s provenance discipline.

DEFECT 3 — Stale README case-index statuses + broken DOI badge (minor)

The README case index showed cases 03/04/06/07/08 as “ In preparation” and 02/05 as “Draft,” though all eight are written and committed. The DOI badge line was malformed markdown (an unterminated image using a `placeholder.co` placeholder URL, not a real Zenodo badge). **Fix:** statuses updated to reflect actual completion; the broken placeholder badge line removed in favor of the working license badge and the real DOI link already present in the citation block.

Per-case-study grading

Each is graded across the five lenses. Scores are justified, not asserted.

01 — The Giza Transmission Hypothesis · Craft 35/35 · Adversarial 96/100

- **Archaeology (A):** Correctly foregrounds the mainstream position (the pyramids are Old Kingdom royal tombs, c. 2600–2500 BCE) and quarantines the “deep-time transmission” idea as an explicitly speculative thought experiment. No fringe claim is smuggled in as fact.
- **Systems (C) / Epistemology (D):** Strongest as a framing piece — it establishes the canonical 11-section structure and the physical-form-survives-but-meaning-may-not thesis that the whole series then tests. This is the series’ constitution.
- **Ethics/theology (E):** The speculative framing is handled with unusual discipline; the reader is never misled about what is evidence and what is hypothesis.
- **Deductions:** It is the most speculative entry, so its adversarial ceiling is inherently lower — a hostile reviewer can dismiss the central hypothesis as unfalsifiable at the margins. The study pre-emptively addresses this with its Falsification Check, which is why it still scores 96 and not lower. Verdict (Speculative/thought-experiment) is correctly labeled.

02 — Aboriginal Oral Traditions · Craft 35/35 · Adversarial 97/100

- **Linguistics/anthropology (B):** Nunn & Reid (2016, Australian Geographer 47(1):11–47) is cited accurately for the sea-level-rise coastal-drowning traditions (~7,000+ years). The claim is stated at the correct strength — genuine cross-millennial information transfer, with the live scholarly debate over exact datability acknowledged rather than buried.
- **Systems (C):** This is the case that legitimately **refines** CIF — it shows active, relational, error-corrected oral transmission outperforming some durable-substrate cases, which pushes back on any naive “durability = fidelity” reading. The verdict (REFINES) is the honest one.
- **Deductions:** Minor — the strength of oral-tradition dating is genuinely contested; the study handles this well but the topic caps the adversarial score just short of the top.

03 — The Antikythera Mechanism · Craft 35/35 · Adversarial 98/100

- **Domain science (B):** 82 fragments, ≥30 surviving gears, Freeth et al. Nature 2006 & 2008 reconstruction work, 21st-century CT/PTM imaging — all reconciled to the record. The “preservation by entropy (the ship sank), not by design” observation is both accurate and analytically sharp.
- **Systems (C):** The Pillar-1-perfect / Pillars-2-5-lost decomposition is the cleanest demonstration in the series that object survival ≠ capability survival. Textbook.
- **Deductions:** Very minor — a hostile reviewer might want explicit acknowledgement that some gear-train interpretations remain debated; the study gestures at this but could name it. Uses the Inheritance-Layer lens (Defect 2), now documented.

04 — The Late Bronze Age Collapse · Craft 35/35 · Adversarial 97/100

- **History (D):** Handles a genuinely contested topic responsibly — presents systems-collapse / multi-causal models (climate, earthquake storms, Sea Peoples, systradecollapse) without overcommitting to the monocausal “Sea Peoples” narrative that popular accounts favor. Linear B survival-by-palace-fire is correctly explained.
- **Systems (C):** The “interconnected systems fail together” analysis is the series’ best treatment of cascading/network failure and correctly maps it to Pillars/Layers 3–5 loss.

- **Deductions:** Longest study (~4,500 w); a hostile reviewer could trim it. Contested causation caps the top score slightly. Inheritance-Layer lens (Defect 2).

05 — Japanese Tsunami Stones • Craft 35/35 • Adversarial 98/100

- **Geoscience (B):** The Aneyoshi stone and the 2011 Tōhoku behavioral evidence are represented accurately. This is the series' cleanest positive result: an inherited boundary marker (Axiom 5) that demonstrably changed behavior generations later.
- **Systems (C) / Ethics (E):** The "warnings must remain warnings after context is lost" point ties directly to Engineering Pillar 5 (non-harm coherence) and is the most actionable lesson in the series.
- **Deductions:** Very minor — the behavioral-causation claim (stones caused villagers to flee to high ground) is strong but partly anecdotal; the study is appropriately measured about it.

06 — Polynesian Wayfinding • Craft 35/35 • Adversarial 97/100

- **Domain science (B) / anthropology:** Hōkūle'a's 1976 Hawai'i→Tahiti voyage (~2,500 mi, ~34 days), Mau Piailug (1932–2010, Satawal) and the revival of non-instrument wayfinding are all accurate. The honest point that the canoe is a **reconstruction** (Pillar 1 recoverable by reverse-engineering, not direct preservation) is exactly right and avoids romanticization.
- **Systems (C):** Best case in the series of a method surviving with **no durable substrate at all** — pure human/relational transmission, near-lost, then revived. Strong test of Axiom 3 (recoverable inheritance) and the Inheritance-Layer method.
- **Deductions:** The 20th-century near-extinction-then-revival means transmission was not continuous; the study is honest that this is reconstruction-assisted continuity, which slightly complicates a clean "survival" reading. Correctly handled.

07 — Göbekli Tepe • Craft 34/35 • Adversarial 95/100

- **Archaeology (A):** Dates (c. 9600–8000 BCE), pillar scale (up to ~5.5 m, ~10–20 t), deliberate burial, and pre-pottery-Neolithic context are accurate. The "iconographic consistency as error-correction across 600+ years" reading is a legitimate and interesting inference.
- **Deductions (why 34, the series' only craft deduction):** This is the study most tempted toward confirmation — several passages assert CIF "working exactly as predicted," which reads slightly as illustration rather than test. The Falsification Check rescues it (it explicitly asks what would contradict CIF and answers honestly), but the surrounding prose leans confirmatory more than the others. A hostile reviewer would also note the "purpose of the monuments" is genuinely unknown, so some Pillar-5 claims are necessarily speculative — the study says so, but could foreground it earlier. Still excellent; just the one entry where the anti-confirmation discipline visibly strains.

08 — Undeciphered Scripts • Craft 35/35 • Adversarial 98/100

- **Linguistics (B):** Indus script, Linear A, and Rongorongo handled precisely — steatite seal durability, Linear A baked-by-palace-fire survival, the 27 surviving Rongorongo tablets, and the correct contrast with Linear B (Ventris 1952) and Egyptian (Champollion 1822–24, Rosetta Stone as the bilingual key). Decipherment status stated accurately (Linear A undeciphered; Indus debated as possibly non-linguistic).
- **Systems (C):** The sharpest statement of the series thesis: "a perfect record is not an inheritance if no one can read it" — Pillars/Layers 1–2 intact, Pillar/Layer 5 severed. The modern-archive analogy (proprietary format, lost key, obsolete software = "a modern Indus seal") is excellent and lands the present-day relevance.

- **Deductions:** Essentially none of substance; the top score is withheld only because the Indus “possibly not a full writing system” debate could be given one more sentence of the opposing view. Inheritance-Layer lens (Defect 2).

Framework document — Craft 33/35 · Adversarial 90/100

- **Strengths:** Genuinely disciplined. The Status/Provenance section handles the UIP→CIF renaming honestly (universality claim failed → “Unified” retired → reframed as a framework), which is the single most important integrity move in the whole repo and is done exactly right. Axioms are explicitly “design propositions, not universal physical laws.” The mandatory Falsification Check is a real anti-confirmation mechanism.
- **Faults (pre-fix):** It is the source of Defect 2 — it defines the Five Engineering Pillars formally but then uses the Five Inheritance Layers (physical form / data / method / intent / semantic meaning) in the canonical Fidelity structure without ever naming or defining them as a set. So the framework doc silently uses two five-part schemes. **Fixed in this commit:** the Inheritance Layers are now defined as a named lens, and the canonical structure states which lens each phase’s Pillar Analysis uses.
- **Post-fix:** Craft rises toward 35; adversarial toward 96.

README — Craft 27/35 · Adversarial 72/100 (pre-fix)

Carried all three defects (wrong axiom list, mislabeled pillars, stale statuses, broken badge). It is the weakest artifact precisely because it is the front door and contradicted the rooms behind it. **Post-fix:** both lists reconciled, statuses corrected, badge fixed → Craft ~34, adversarial ~95.

Series-level assessment

What the series does exceptionally well (holds at 35-craft standard):

1. **Anti-confirmation by design.** Every study can lose. The taxonomy includes CHALLENGED and FALSIFIED AS FRAMED, case 02 refines the framework, and no study is allowed to be confirmation-only. This is rare and is the series’ single greatest strength.
2. **Fact discipline.** Every checkable specific reconciled to the mainstream record. No fabricated citations, no invented percentages — fidelity is assessed qualitatively, as promised.
3. **Honest provenance.** The UIP→CIF history is stated plainly everywhere it matters. No “first/proven/universal/patented” overclaiming.
4. **Cumulative argument.** The eight cases genuinely triangulate one thesis — physical/data survival ≠ meaning/capability survival — from positive (05), negative (03, 08), collapse (04), no-substrate (06), pre-writing (07), oral (02), and speculative (01) angles.

What kept the series off 100 (pre-fix): the two internal contradictions and the stale README — all now corrected. The deepest honest finding is that the anti-drift framework drifted its own vocabulary; that is recorded rather than hidden.

Remaining, disclosed limitations (not defects — inherent):

- The synthesis capstone is intentionally unwritten (`synthesis/README.md` is a stub) — correct sequencing (pattern read out of the cases, not imposed), but the series is not yet “complete” until it exists.

- A series-level DOI is not yet minted; the current DOI resolves to the Giza foundational study only. The README states this honestly.
- Case 07 is the one place the anti-confirmation discipline visibly strains.

Post-fix series verdict: 35/35 craft · 97/100 adversarial. The three points withheld from 100 reflect the inherent limits above (unwritten synthesis, single-study DOI, case-07 strain) — not concealment or error. Under the honesty covenant, that is the truthful ceiling for the series as it stands today, and it is a genuinely high one.

Reviewed under the Remnant Fieldworks highest-honesty quality bar. Faults led with; scores justified; nothing softened.